

Analysis of Liver Hepatocellular Carcinoma Data in R

Bo Yang, Rui Nie, Yueying Hu

Introduction

Hepatocellular carcinoma (LIHC/HCC) is the one of the leading cause of cancer death, which accounts for 8.2% of global cancer incidence and 4.7% of cancer-related mortality [1]. Previous study suggested that lifestyles and personal habits, including drinking alcohol, smoking, are positively associated with the onset LIHC [2], which has greatly enhanced the efficiency of screening for the disease by health providers. Nevertheless, the difficulty for disease detection remains alarmingly high. Therefore, an indicator with higher sensitivity would advance the diagnosis stage and prolong patients' survival.

RNA Sequencing (RNA-seq) using the capabilities of high-throughput sequencing methods provides insight into the transcriptome of a cell. Through understanding the gene expression patterns reflected by RNA sequencing results, we aim to locate indicator genes for LIHC with the help of statistical learning methods. Furthermore, given the massive amount of existing genes in human genome, our problem can be formulated as a high-dimensional problem. One of our side goal is to perform variable selection to reduce the risk of overfitting.

Methods

LIHC data from TCGA

We're primarily interested in analyzing RNA-seq data provided through The Cancer Genome Atlas (TCGA) for liver cancer patients. The dataset includes data from 424 LIHC participants and over 20,000 genes for expression level analysis. The inclusion of both normal and tumor samples allows us for comparative analysis between cancerous and non-cancerous liver tissues.

To download and prepare the LIHC data from TCGA biolinks, we used `GDCquery`, `GDCdownload`, and `GDCprepare` functions for data accessing.

```
query_TCGA_LIHC = GDCquery(  
  project = "TCGA-LIHC",  
  data.category = "Transcriptome Profiling",  
  experimental.strategy = "RNA-Seq",  
  workflow.type = "STAR - Counts",  
  data.type = 'Gene Expression Quantification')  
  
GDCdownload(query = query_TCGA_LIHC)  
  
tcga_data_LIHC = GDCprepare(query_TCGA_LIHC)
```

```
## | | 0% |
```

```
# transform a large summarized experiment into readable matrix
count=assay(tcga_data_LIHC)
```

Normalization for reads count help account for gene length and sequencing depth. Fragments Per Kilobase of transcript per Million mapped reads (FPKM) provides a comparable measure to interpret and utilize RNA-seq results. We standardized our data into FPKM format for fair representation and comparison gene expressions.

```
## load gene annotation data
file.annotations <- system.file("extdata",
                                "Biomart.annotations.hg38.txt",
                                package="countToFPKM")
gene.annotations <- read.table(file.annotations, sep="\t", header=TRUE)

## combine gene annotation and expression data
genelength=data.frame(gene=gene.annotations$Gene_ID,
                      length=gene.annotations$length,
                      type=gene.annotations$Gene_type)

rownames(count)<-substring(rownames(count),1,15)
df_temp = merge(genelength, count, by.x="gene",by.y="row.names")

## FPKM calculation by hand
countToFpkm <- function(counts, effLen)
{
  N <- sum(counts)
  exp( log(counts) + log(1e9) - log(effLen) - log(N) )
}

df<-countToFpkm(df_temp[, -c(1:3)], df_temp$length)
rownames(df)<-df_temp$gene; X = as.matrix(t(df));dim(df)
```

```
## [1] 20063 424
```

Preprocessing

We collected the sample IDs from normal samples and from tumor samples separately and retrieved their index for future train-test dataset split. In summary, out of 424 samples, there are 50 normal samples and 374 tumor samples. This suggests we consider additional metrics such as precision, F1, or AUC-ROC scores to account for potential bias introduced by the imbalanced labels aside from crude accuracy for classification tasks.

```
## retrieve the label for each sample
sample_ID = colnames(df)
tumor_sample_ID = sample_ID[which(substring(sample_ID,14,15)!=11)]
normal_sample_ID = sample_ID[which(substring(sample_ID,14,15)==11)]

## index for each label group
tumor_index = which(sample_ID %in% tumor_sample_ID)
normal_index = which(sample_ID %in% normal_sample_ID)
label = rep(0, length(sample_ID)); label[tumor_index] = 1
label = factor(label, labels = c("normal", "tumor")); table(label)
```

```
## label
## normal  tumor
##      50    374
```

Variable Selection

Of all the variables (genes) selected for our preprocessed data, we'd like to filter out the ones with constant values across samples, as they provide no information in differentiating tumor samples from normal samples. This leads to the removal of 381 genes from the entire sample.

```
sd_x = apply(X, 2, sd)
ind = which(sd_x == 0)
X = X[, -ind]; length(ind)
```


```
## [1] 381
```

For dimension reduction, we visualized samples with different labels through principle components analysis, which linearly combined all input variables to form principle component according to their variance in a descending order. We plotted all samples according to their first and second principle component, which explained a total of 24.89% variance across all genes. The clustering of tumor samples indicates that tumor samples may have a shared gene expression pattern.

```
res.pca = PCA(X, scale.unit = TRUE, graph = FALSE)
eigenvalue = get_eigenvalue(res.pca); head(eigenvalue[,3], 2)
```

```
##      Dim.1      Dim.2
## 19.76191 24.89457
```

```
fviz_pca_ind(res.pca, axes = c(1, 2), col.ind = label, label = FALSE)
```



A dataset having as many as 19682 covariates versus only 424 samples can be regarded as a high-dimensional problem and have a strong tendency to cause overfit in prediction models.

To select variables while maintaining interpretation along the process, we will use the variance filtering method to select genes with a more variable expression activity from the samples we included. We kept the variance cutoff to be as high as 0.95, leaving only 5% (1004) of the original genes for future statistical analysis.

```
Xnew = t(varFilter(t(X), var.func=IQR, var.cutoff = 0.95, filterByQuantile = TRUE));dim(Xnew)
```

```
## [1] 424 985
```

Statistical Learning

For building predictive statistical learning methods in LHC classification, we continue with 985 selected genes using Elastic Net, decision tree, and random forest and assess their performances under the idea of train-test split validation. We will also use cross-validation to tune the hyperparameters for each model and compare their performances in terms of accuracy, precision, recall, F1 score, and AUC-ROC score.

Results

Train-Test Split

Considering the imbalanced labels of two classes, it is important to preserve the representation of both classes in the training and testing samples to ensure the model is not overly biasing towards one class over the other. As a consequence, we use `createDataPartition` to create a stratified split for training and testing purposes with a 3:1 ratio.

```
## statistical learning
set.seed(123)
train_index = createDataPartition(label, p = 0.75, list = FALSE)
trainX = Xnew[train_index, ]; testX = Xnew[-train_index, ]
train_label = label[train_index]; test_label = label[-train_index]
```

Elastic Net

Elastic Net is a penalized regression method that combines the L1 and L2 penalties of Lasso and Ridge regression. The L1 penalty suppress the coefficients of some variables to zero, which is equivalent to variable selection. The L2 penalty shrinks the coefficients of some variables towards a smaller scale. We take the idea of logistic regression and use below formula as the loss function for Elastic Net in classification problems:

$$\min_{\beta_0, \beta} \frac{1}{N} \sum_{i=1}^N \log(1 + \exp(-y_i(\beta_0 + x_i^T \beta))) + \lambda \left[\frac{1}{2}(1 - \alpha) \sum_{j=1}^p \beta_j^2 + \alpha \sum_{j=1}^p |\beta_j| \right]$$

We determine the alpha and lambda coefficient through grid search, where alpha determines the balance between L1 and L2 penalty and lambda determines the strength of the penalty.

```
## grid search for hyperparameter tuning
set.seed(123)
alpha_grid = seq(0, 1, 0.1)
lambda_grid = seq(0.001, 0.1, 0.001)
cv_control <- trainControl(method = "cv", number = 5, classProbs = TRUE)
elastic <- train(x = trainX, y = train_label, method = "glmnet",
  trControl = cv_control, family = "binomial", tuneGrid = expand.grid(alpha = alpha_grid, lambda = lambda_grid))
elastic$bestTune
```

```
##      alpha lambda
## 236    0.2  0.036
```

```
## Evaluate the performance of Elastic Net
```

```
set.seed(123)
```

```
elastic_model = glmnet(trainX, train_label, family="binomial", alpha = 0.2, lambda = 0.036)
```

```
elastic_pred = predict(elastic_model, testX, type = "class")
```

```
result.elastic = confusionMatrix(factor(elastic_pred), test_label)$byClass[c(1, 4, 5, 6, 7)]
```

```
## Variables selected by Elastic Net
```

```
coef.elastic = coef(elastic_model, s = "lambda.min")
```

```
coef.elastic = coef.elastic[coef.elastic[,1]!=0,];head(coef.elastic, 5)
```

```
##      (Intercept) ENSG00000008517 ENSG00000011304 ENSG00000044574 ENSG00000064601
##      2.5789832      0.2045341      2.5925027      0.5199208      0.2184742
```

```
length(coef.elastic[-1])
```

```
## [1] 120
```